



In [41]:

```
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import sqlite3
import csv
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from wordcloud import WordCloud
import re
import os
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
import numpy as np
from sqlalchemy import create_engine # database connection
import datetime as dt
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem.snowball import SnowballStemmer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.multiclass import OneVsRestClassifier
from sklearn.linear_model import SGDClassifier
from sklearn import metrics
from sklearn.metrics import f1_score, precision_score, recall_score
from sklearn import svm
from sklearn.linear_model import LogisticRegression
from skmultilearn.adapt import mlknn

from skmultilearn.problem_transform import ClassifierChain
from skmultilearn.problem_transform import BinaryRelevance
from skmultilearn.problem_transform import LabelPowerset
from sklearn.naive_bayes import GaussianNB
from datetime import datetime
import pickle
from sklearn.externals import joblib
```

Stack Overflow: Tag Prediction

1. Business Problem

1.1 Description

Description

Stack Overflow is the largest, most trusted online community for developers to learn, share their programming knowledge, and build their careers.

Stack Overflow is something which every programmer use one way or another. Each month, over 50 million developers come to Stack Overflow to learn, share their knowledge, and build their careers. It features questions and answers on a wide range of topics in computer programming. The website serves as a platform for users to ask and answer questions, and, through membership and active participation, to vote questions and

answers up or down and edit questions and answers in a fashion similar to a wiki or Digg. As of April 2014 Stack Overflow has over 4,000,000 registered users, and it exceeded 10,000,000 questions in late August 2015. Based on the type of tags assigned to questions, the top eight most discussed topics on the site are: Java, JavaScript, C#, PHP, Android, jQuery, Python and HTML.

Problem Statement

Suggest the tags based on the content that was there in the question posted on Stackoverflow.

Source: <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/>

1.2 Source / useful links

Data Source : <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>

(<https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>)

Youtube : <https://youtu.be/nNDqbUhtIRg> (<https://youtu.be/nNDqbUhtIRg>)

Research paper : <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/tagging-1.pdf>

(<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/tagging-1.pdf>)

Research paper : <https://dl.acm.org/citation.cfm?id=2660970&dl=ACM&coll=DL> (<https://dl.acm.org/citation.cfm?id=2660970&dl=ACM&coll=DL>)

1.3 Real World / Business Objectives and Constraints

1. Predict as many tags as possible with high precision and recall.
2. Incorrect tags could impact customer experience on StackOverflow.
3. No strict latency constraints.

2. Machine Learning problem

2.1 Data

2.1.1 Data Overview

Refer: <https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>

(<https://www.kaggle.com/c/facebook-recruiting-iii-keyword-extraction/data>)

All of the data is in 2 files: Train and Test.

Train.csv contains 4 columns: Id,Title,Body,Tags.

Test.csv contains the same columns but without the Tags, which you are to predict.

Size of Train.csv - 6.75GB

Size of Test.csv - 2GB

Number of rows in Train.csv = 6034195

The questions are randomized and contains a mix of verbose text sites as well as sites related to math and programming. The number of questions from each site may vary, and no filtering has been performed on the questions (such as closed questions).

Data Field Explanation

Dataset contains 6,034,195 rows. The columns in the table are:

Id - Unique identifier for each question

Title - The question's title

Body - The body of the question

Tags - The tags associated with the question in a space-separated format (all lowercase, should not contain tabs '\t' or ampersands '&')

2.1.2 Example Data point

Title: Implementing Boundary Value Analysis of Software Testing in a C++ program?

Body :

```

#include<
iostream>\n
#include<
stdlib.h>\n\n
using namespace std;\n\n
int main()\n
{\n
    int n,a[n],x,c,u[n],m[n],e[n][4];\n
    cout<<"Enter the number of variables";\n            ci
n>>n;\n\n
    cout<<"Enter the Lower, and Upper Limits of the var
iables";\n
    for(int y=1; y<n+1; y++)\n
    {\n
        cin>>m[y];\n
        cin>>u[y];\n
    }\n
    for(x=1; x<n+1; x++)\n
    {\n
        a[x] = (m[x] + u[x])/2;\n
    }\n
    c=(n*4)-4;\n
    for(int a1=1; a1<n+1; a1++)\n
    {\n\n
        e[a1][0] = m[a1];\n
        e[a1][1] = m[a1]+1;\n
        e[a1][2] = u[a1]-1;\n
        e[a1][3] = u[a1];\n
    }\n
    for(int i=1; i<n+1; i++)\n
    {\n
        for(int l=1; l<=i; l++)\n
        {\n
            if(l!=1)\n
            {\n
                cout<<a[l]<<"\\t";\n
            }\n
        }\n
        for(int j=0; j<4; j++)\n
        {\n
            cout<<e[i][j];\n
            for(int k=0; k<n-(i+1); k++)\n
            {\n
                cout<<a[k]<<"\\t";\n
            }\n
            cout<<"\\n";\n
        }\n
    }\n
    \n\n
    system("PAUSE");\n
    return 0;    \n

```

```
} \n
```

```
\n \n
```

<p>The answer should come in the form of a table like</p>\n \n

```
<pre><code>
```

```
1          50          50 \n
2          50          50 \n
99         50          50 \n
100        50          50 \n
50         1           50 \n
50         2           50 \n
50        99          50 \n
50       100          50 \n
50        50          1 \n
50        50          2 \n
50        50          99 \n
50        50         100 \n
```

```
</code></pre>\n \n
```

<p>if the no of inputs is 3 and their ranges are \n

```
1,100 \n
```

```
1,100 \n
```

```
1,100 \n
```

(could be varied too)</p>\n \n

<p>The output is not coming, can anyone correct the code or tell me what's wrong?</p>\n'

Tags : 'c++ c'

2.2 Mapping the real-world problem to a Machine Learning Problem

2.2.1 Type of Machine Learning Problem

It is a multi-label classification problem

Multi-label Classification: Multilabel classification assigns to each sample a set of target labels. This can be thought as predicting properties of a data-point that are not mutually exclusive, such as topics that are relevant for a document. A question on Stackoverflow might be about any of C, Pointers, FileIO and/or memory-management at the same time or none of these.

___Credit___: <http://scikit-learn.org/stable/modules/multiclass.html>

2.2.2 Performance metric

Micro-Averaged F1-Score (Mean F Score) : The F1 score can be interpreted as a weighted average of the precision and recall, where an F1 score reaches its best value at 1 and worst score at 0. The relative contribution of precision and recall to the F1 score are equal. The formula for the F1 score is:

$$F1 = 2 * (precision * recall) / (precision + recall)$$

In the multi-class and multi-label case, this is the weighted average of the F1 score of each class.

'Micro f1 score':

Calculate metrics globally by counting the total true positives, false negatives and false positives. This is a better metric when we have class imbalance.

'Macro f1 score':

Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.

<https://www.kaggle.com/wiki/MeanFScore> (<https://www.kaggle.com/wiki/MeanFScore>)
http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html (http://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html)

Hamming loss : The Hamming loss is the fraction of labels that are incorrectly predicted.
<https://www.kaggle.com/wiki/HammingLoss> (<https://www.kaggle.com/wiki/HammingLoss>)

3. Exploratory Data Analysis

3.1 Data Loading and Cleaning

3.1.1 Using Pandas with SQLite to Load the data

In [0]:

```
#Creating db file from csv
#Learn SQL: https://www.w3schools.com/sql/default.asp
if not os.path.isfile('train.db'):
    start = datetime.now()
    disk_engine = create_engine('sqlite:///train.db')
    start = dt.datetime.now()
    chunksize = 180000
    j = 0
    index_start = 1
    for df in pd.read_csv('Train.csv', names=['Id', 'Title', 'Body', 'Tags'], chunksize=chunksize):
        df.index += index_start
        j+=1
        print('{} rows'.format(j*chunksize))
        df.to_sql('data', disk_engine, if_exists='append')
        index_start = df.index[-1] + 1
    print("Time taken to run this cell :", datetime.now() - start)
```

3.1.2 Counting the number of rows

In [0]:

```

if os.path.isfile('train.db'):
    start = datetime.now()
    con = sqlite3.connect('train.db')
    num_rows = pd.read_sql_query("""SELECT count(*) FROM data""", con)
    #Always remember to close the database
    print("Number of rows in the database :", "\n", num_rows['count(*)'].values[0])
    con.close()
    print("Time taken to count the number of rows :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the above cell to ge

```

Number of rows in the database :

6034196

Time taken to count the number of rows : 0:01:15.750352

3.1.3 Checking for duplicates

In [0]:

```

#Learn SQL: https://www.w3schools.com/sql/default.asp
if os.path.isfile('train.db'):
    start = datetime.now()
    con = sqlite3.connect('train.db')
    df_no_dup = pd.read_sql_query('SELECT Title, Body, Tags, COUNT(*) as cnt_dup FR
    con.close()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the first to generat

```

Time taken to run this cell : 0:04:33.560122

In [0]:

```

df_no_dup.head()
# we can observe that there are duplicates

```

Out[6]:

	Title	Body	Tags	cnt_dup
0	Implementing Boundary Value Analysis of S...	<pre>#include<iosstream>\n#include&...	c++ c	1
1	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamical...	c# silverlight data-binding	1
2	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamical...	c# silverlight data-binding columns	1
3	java.lang.NoClassDefFoundError: javax/serv...	<p>I followed the guide in <a href="http://sta...	jsp jstl	1
4	java.sql.SQLException: [Microsoft] [ODBC Dri...	<p>I use the following code</p>\n\n<pre>#include<iosstream>\n#include&...	java jdbc	2

In [0]:

```
print("number of duplicate questions :", num_rows['count(*)'].values[0] - df_no_dup.
```

```
number of duplicate questions : 1827881 ( 30.2920389063 % )
```

In [0]:

```
# number of times each question appeared in our database
df_no_dup.cnt_dup.value_counts()
```

Out[8]:

```
1    2656284
2    1272336
3     277575
4         90
5         25
6          5
Name: cnt_dup, dtype: int64
```

In [0]:

```
start = datetime.now()
df_no_dup["tag_count"] = df_no_dup["Tags"].apply(lambda text: len(text.split(" ")))
# adding a new feature number of tags per question
print("Time taken to run this cell :", datetime.now() - start)
df_no_dup.head()
```

```
Time taken to run this cell : 0:00:03.169523
```

Out[9]:

	Title	Body	Tags	cnt_dup
0	Implementing Boundary Value Analysis of S...	<pre> <code>#include<i>iostream</i>>\n#include&...	c++ c	1
1	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...	c# silverlight data-binding	1
2	Dynamic Datagrid Binding in Silverlight?	<p>I should do binding for datagrid dynamicall...	c# silverlight data-binding columns	1
3	java.lang.NoClassDefFoundError: javax/serv...	<p>I followed the guide in <a href="http://sta...	jsp jstl	1
4	java.sql.SQLException:[Microsoft][ODBC Dri...	<p>I use the following code</p>\n\n<pre> <code>...	java jdbc	2

In [0]:

```
# distribution of number of tags per question
df_no_dup.tag_count.value_counts()
```

Out[10]:

```
3    1206157
2    1111706
4     814996
1     568298
5     505158
Name: tag_count, dtype: int64
```

In [0]:

```
#Creating a new database with no duplicates
if not os.path.isfile(path+'train_no_dup.db'):
    disk_dup = create_engine("sqlite:///train_no_dup.db")
    no_dup = pd.DataFrame(df_no_dup, columns=['Title', 'Body', 'Tags'])
    no_dup.to_sql('no_dup_train',disk_dup)
```

In [10]:

```
#This method seems more appropriate to work with this much data.
#creating the connection with database file.
if os.path.isfile('train_no_dup.db'):
    start = datetime.now()
    con = sqlite3.connect('train_no_dup.db')
    tag_data = pd.read_sql_query("""SELECT Tags FROM no_dup_train""", con)
    #Always remember to close the database
    con.close()

    # Let's now drop unwanted column.
    tag_data.drop(tag_data.index[0], inplace=True)
    #Printing first 5 columns from our data frame
    tag_data.head()
    print("Time taken to run this cell :", datetime.now() - start)
else:
    print("Please download the train.db file from drive or run the above cells to g
```

Time taken to run this cell : 0:01:46.493856

3.2 Analysis of Tags

3.2.1 Total number of unique tags

In [0]:

```
# Importing & Initializing the "CountVectorizer" object, which
#is scikit-learn's bag of words tool.

#by default 'split()' will tokenize each tag using space.
vectorizer = CountVectorizer(tokenizer = lambda x: x.split())
# fit_transform() does two functions: First, it fits the model
# and learns the vocabulary; second, it transforms our training data
# into feature vectors. The input to fit_transform should be a list of strings.
tag_dtm = vectorizer.fit_transform(tag_data['Tags'])
```

In [12]:

```
print("Number of data points :", tag_dtm.shape[0])
print("Number of unique tags :", tag_dtm.shape[1])
```

Number of data points : 4206314
Number of unique tags : 42048

In [13]:

```
#'get_feature_name()' gives us the vocabulary.
tags = vectorizer.get_feature_names()
#Lets look at the tags we have.
print("Some of the tags we have :", tags[:10])
```

Some of the tags we have : ['.a', '.app', '.asp.net-mvc', '.aspxaut
h', '.bash-profile', '.class-file', '.cs-file', '.doc', '.drv', '.ds-
store']

3.2.3 Number of times a tag appeared

In [0]:

```
# https://stackoverflow.com/questions/15115765/how-to-access-sparse-matrix-elements
#Lets now store the document term matrix in a dictionary.
freqs = tag_dtm.sum(axis=0).A1
result = dict(zip(tags, freqs))
```

In [17]:

```
#Saving this dictionary to csv files.
if not os.path.isfile('tag_counts_dict_dtm.csv'):
    with open('tag_counts_dict_dtm.csv', 'w') as csv_file:
        writer = csv.writer(csv_file)
        for key, value in result.items():
            writer.writerow([key, value])
tag_df = pd.read_csv("tag_counts_dict_dtm.csv", names=['Tags', 'Counts'])
tag_df.head()
```

Out[17]:


	Tags	Counts
0	.a	18
1	.app	37
2	.asp.net-mvc	1
3	.aspxauth	21
4	.bash-profile	138

In [0]:

```
tag_df_sorted = tag_df.sort_values(['Counts'], ascending=False)
tag_counts = tag_df_sorted['Counts'].values
```


In [19]:

```
plt.plot(tag_counts)
plt.title("Distribution of number of times tag appeared questions")
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
```



In [20]:

```
plt.plot(tag_counts[0:10000])
plt.title('first 10k tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:10000:25]), tag_counts[0:10000:25])
```




400	[331505	44829	22429	17728	13364	11162	10029	9148	8054	
7151										
	6466	5865	5370	4983	4526	4281	4144	3929	3750	359
3										
	3453	3299	3123	2989	2891	2738	2647	2527	2431	233
1										
	2259	2186	2097	2020	1959	1900	1828	1770	1723	167
3										
	1631	1574	1532	1479	1448	1406	1365	1328	1300	126
6										
	1245	1222	1197	1181	1158	1139	1121	1101	1076	105
6										
	1038	1023	1006	983	966	952	938	926	911	89
1										
	882	869	856	841	830	816	804	789	779	77
0										
	752	743	733	725	712	702	688	678	671	65
8										
	650	643	634	627	616	607	598	589	583	57
7										
	568	559	552	545	540	533	526	518	512	50
6										
	500	495	490	485	480	477	469	465	457	45
0										
	447	442	437	432	426	422	418	413	408	40
3										

5	398	393	388	385	381	378	374	370	367	36
2	361	357	354	350	347	344	342	339	336	33
1	330	326	323	319	315	312	309	307	304	30
6	299	296	293	291	289	286	284	281	278	27
4	275	272	270	268	265	262	260	258	256	25
6	252	250	249	247	245	243	241	239	238	23
9	234	233	232	230	228	226	224	222	220	21
3	217	215	214	212	210	209	207	205	204	20
9	201	200	199	198	196	194	193	192	191	18
7	188	186	185	183	182	181	180	179	178	17
5	175	174	172	171	170	169	168	167	166	16
5	164	162	161	160	159	158	157	156	156	15
6	154	153	152	151	150	149	149	148	147	14
7	145	144	143	142	142	141	140	139	138	13
0	137	136	135	134	134	133	132	131	130	13
3	129	128	128	127	126	126	125	124	124	12
7	123	122	122	121	120	120	119	118	118	11
1	117	116	116	115	115	114	113	113	112	11
6	111	110	109	109	108	108	107	106	106	10
1	105	105	104	104	103	103	102	102	101	10
6	100	100	99	99	98	98	97	97	96	9
1	95	95	94	94	93	93	93	92	92	9
6	91	90	90	89	89	88	88	87	87	8
2	86	86	85	85	84	84	83	83	83	8
8	82	82	81	81	80	80	80	79	79	7
5	78	78	78	77	77	76	76	76	75	7
2]	75	74	74	74	73	73	73	73	72	7

In [21]:

```
plt.plot(tag_counts[0:1000])
plt.title('first 1k tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:1000:5]), tag_counts[0:1000:5])
```



```
200 [331505 221533 122769 95160 62023 44829 37170 31897 26925
24537
 22429 21820 20957 19758 18905 17728 15533 15097 14884 1370
3
 13364 13157 12407 11658 11228 11162 10863 10600 10350 1022
4
 10029 9884 9719 9411 9252 9148 9040 8617 8361 816
3
 8054 7867 7702 7564 7274 7151 7052 6847 6656 655
3
 6466 6291 6183 6093 5971 5865 5760 5577 5490 541
1
 5370 5283 5207 5107 5066 4983 4891 4785 4658 454
9
 4526 4487 4429 4335 4310 4281 4239 4228 4195 415
9
 4144 4088 4050 4002 3957 3929 3874 3849 3818 379
7
 3750 3703 3685 3658 3615 3593 3564 3521 3505 348
3
 3453 3427 3396 3363 3326 3299 3272 3232 3196 316
8
 3123 3094 3073 3050 3012 2989 2984 2953 2934 290
3
 2891 2844 2819 2784 2754 2738 2726 2708 2681 266
9
 2647 2621 2604 2594 2556 2527 2510 2482 2460 244
4
```

```

2431 2409 2395 2380 2363 2331 2312 2297 2290 228
1
2259 2246 2222 2211 2198 2186 2162 2142 2132 210
7
2097 2078 2057 2045 2036 2020 2011 1994 1971 196
5
1959 1952 1940 1932 1912 1900 1879 1865 1855 184
1
1828 1821 1813 1801 1782 1770 1760 1747 1741 173
4
1723 1707 1697 1688 1683 1673 1665 1656 1646 163
9]


```

In [22]:

```

plt.plot(tag_counts[0:500])
plt.title('first 500 tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.show()
print(len(tag_counts[0:500:5]), tag_counts[0:500:5])

```



```

100 [331505 221533 122769 95160 62023 44829 37170 31897 26925
24537
22429 21820 20957 19758 18905 17728 15533 15097 14884 1370
3
13364 13157 12407 11658 11228 11162 10863 10600 10350 1022
4
10029 9884 9719 9411 9252 9148 9040 8617 8361 816
3
8054 7867 7702 7564 7274 7151 7052 6847 6656 655
3
6466 6291 6183 6093 5971 5865 5760 5577 5490 541
1
5370 5283 5207 5107 5066 4983 4891 4785 4658 454
9
4526 4487 4429 4335 4310 4281 4239 4228 4195 415
9
4144 4088 4050 4002 3957 3929 3874 3849 3818 379
7
3750 3703 3685 3658 3615 3593 3564 3521 3505 348
3]


```


In [23]:

```
plt.plot(tag_counts[0:100], c='b')
plt.scatter(x=list(range(0,100,5)), y=tag_counts[0:100:5], c='orange', label="quantiles with 0.25 difference")
plt.scatter(x=list(range(0,100,25)), y=tag_counts[0:100:25], c='m', label = "quantiles with 0.05 difference")

for x,y in zip(list(range(0,100,25)), tag_counts[0:100:25]):
    plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500))

plt.title('first 100 tags: Distribution of number of times tag appeared questions')
plt.grid()
plt.xlabel("Tag number")
plt.ylabel("Number of times tag appeared")
plt.legend()
plt.show()
print(len(tag_counts[0:100:5]), tag_counts[0:100:5])
```



```
20 [331505 221533 122769 95160 62023 44829 37170 31897 26925 2
4537
22429 21820 20957 19758 18905 17728 15533 15097 14884 1370
3]
```

In [24]:

```
# Store tags greater than 10K in one list
lst_tags_gt_10k = tag_df[tag_df.Counts>10000].Tags
#Print the length of the list
print ('{} Tags are used more than 10000 times'.format(len(lst_tags_gt_10k)))
# Store tags greater than 100K in one list
lst_tags_gt_100k = tag_df[tag_df.Counts>100000].Tags
#Print the length of the list.
print ('{} Tags are used more than 100000 times'.format(len(lst_tags_gt_100k)))
```

```
153 Tags are used more than 10000 times
14 Tags are used more than 100000 times
```

Observations:

1. There are total 153 tags which are used more than 10000 times.
2. 14 tags are used more than 100000 times.
3. Most frequent tag (i.e. c#) is used 331505 times.

4. Since some tags occur much more frequently than others, Micro-averaged F1-score is the appropriate metric for this problem.

3.2.4 Tags Per Question

In [25]:

```
#Storing the count of tag in each question in list 'tag_count'
tag_quest_count = tag_dtm.sum(axis=1).tolist()
#Converting list of lists into single list, we will get [[3], [4], [2], [2], [3]] a
tag_quest_count=[int(j) for i in tag_quest_count for j in i]
print ('We have total {} datapoints.'.format(len(tag_quest_count)))

print(tag_quest_count[:5])
```

We have total 4206314 datapoints.
[3, 4, 2, 2, 3]


In [26]:

```
print( "Maximum number of tags per question: %d"%max(tag_quest_count))
print( "Minimum number of tags per question: %d"%min(tag_quest_count))
print( "Avg. number of tags per question: %f"% ((sum(tag_quest_count)*1.0)/len(tag_
```

Maximum number of tags per question: 5
Minimum number of tags per question: 1
Avg. number of tags per question: 2.899440

In [27]:

```
sns.countplot(tag_quest_count, palette='gist_rainbow')
plt.title("Number of tags in the questions ")
plt.xlabel("Number of Tags")
plt.ylabel("Number of questions")
plt.show()
```



Observations:

1. Maximum number of tags per question: 5
2. Minimum number of tags per question: 1

3. Avg. number of tags per question: 2.899
4. Most of the questions are having 2 or 3 tags

3.2.5 Most Frequent Tags


In [28]:

```
# Plotting word cloud
start = datetime.now()

# Lets first convert the 'result' dictionary to 'list of tuples'
tup = dict(result.items())

#Initializing WordCloud using frequencies of tags.
wordcloud = WordCloud(    background_color='black',
                           width=1600,
                           height=800,
                           ).generate_from_frequencies(tup)

fig = plt.figure(figsize=(30,20))
plt.imshow(wordcloud)
plt.axis('off')
plt.tight_layout(pad=0)
fig.savefig("tag.png")
plt.show()
print("Time taken to run this cell :", datetime.now() - start)
```



Time taken to run this cell : 0:00:04.645941


Observations:

A look at the word cloud shows that "c#", "java", "php", "asp.net", "javascript", "c++" are some of the most frequent tags.

3.2.6 The top 20 tags

In [29]:

```
i=np.arange(30)
tag_df_sorted.head(30).plot(kind='bar')
plt.title('Frequency of top 20 tags')
plt.xticks(i, tag_df_sorted['Tags'])
plt.xlabel('Tags')
plt.ylabel('Counts')
plt.show()
```



In [32]:

```
tag_df_sorted['Tags']
```

Out[32]:

```
4337      c#
18069     java
27249     php
18157  javascript
1234    android
...
29936  rbindlist
29934    rbga
29930    rbar
2925   azureus
42047  zzt-oop
Name: Tags, Length: 42048, dtype: object
```

Observations:

1. Majority of the most frequent tags are programming language.
2. C# is the top most frequent programming language.
3. Android, IOS, Linux and windows are among the top most frequent operating systems.

3.3 Cleaning and preprocessing of Questions

3.3.1 Preprocessing

1. Sample 1M data points
2. Separate out code-snippets from Body
3. Remove Special characters from Question title and description (not in code)
4. Remove stop words (Except 'C')
5. Remove HTML Tags
6. Convert all the characters into small letters
7. Use SnowballStemmer to stem the words

In [7]:

```
def striphtml(data):  
    cleanr = re.compile('<.*?>')  
    cleantext = re.sub(cleanr, ' ', str(data))  
    return cleantext  
stop_words = set(stopwords.words('english'))  
stemmer = SnowballStemmer("english")
```

In [8]:

```

#http://www.sqlitetutorial.net/sqlite-python/create-tables/
def create_connection(db_file):
    """ create a database connection to the SQLite database
        specified by db_file
    :param db_file: database file
    :return: Connection object or None
    """
    try:
        conn = sqlite3.connect(db_file)
        return conn
    except Error as e:
        print(e)

    return None

def create_table(conn, create_table_sql):
    """ create a table from the create_table_sql statement
    :param conn: Connection object
    :param create_table_sql: a CREATE TABLE statement
    :return:
    """
    try:
        c = conn.cursor()
        c.execute(create_table_sql)
    except Error as e:
        print(e)

def checkTableExists(dbcon):
    cursr = dbcon.cursor()
    str = "select name from sqlite_master where type='table'"
    table_names = cursr.execute(str)
    print("Tables in the databse:")
    tables = table_names.fetchall()
    print(tables[0][0])
    return(len(tables))

def create_database_table(database, query):
    conn = create_connection(database)
    if conn is not None:
        create_table(conn, query)
        checkTableExists(conn)
    else:
        print("Error! cannot create the database connection.")
    conn.close()

sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question text
create_database_table("Processed.db", sql_create_table)

```

Tables in the databse:
QuestionsProcessed

In [39]:

```
# http://www.sqlitetutorial.net/sqlite-delete/
# https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table
start = datetime.now()
read_db = 'train_no_dup.db'
write_db = 'Processed.db'
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY RANDOM()")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")
print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:23:09.457459

we create a new data base to store the sampled and preprocessed questions

In [0]:

```
#http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/

start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0
questions_proccesed = 0
for row in reader:

    is_code = 0

    title, question, tags = row[0], row[1], row[2]

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTAL
    question=striphtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    question=str(title)+" "+str(question)
    question=re.sub(r'[^A-Za-z]+',' ',question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question exceptt for the let
    question=' '.join(str(stemmer.stem(j)) for j in words if j not in stop_words an

    len_post+=len(question)
    tup = (question,code,tags,x,len(question),is_code)
    questions_proccesed += 1
    writer.execute("insert into QuestionsProcessed(question,code,tags,words_pre,words_post) values(?,?,?,?,?)")
    if (questions_proccesed%100000==0):
        print("number of questions completed=",questions_proccesed)

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_dup_avg_len_pre)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_dup_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code*100.0)/questions_proccesed))

print("Time taken to run this cell :", datetime.now() - start)
```

```
number of questions completed= 100000
number of questions completed= 200000
number of questions completed= 300000
number of questions completed= 400000
number of questions completed= 500000
number of questions completed= 600000
number of questions completed= 700000
```



```
number of questions completed= 800000  
number of questions completed= 900000  
Avg. length of questions(Title+Body) before processing: 1169  
Avg. length of questions(Title+Body) after processing: 327  
Percent of questions containing code: 57  
Time taken to run this cell : 0:47:05.946582
```

In [0]:

```
# dont forget to close the connections, or else you will end up with locks  
conn_r.commit()  
conn_w.commit()  
conn_r.close()  
conn_w.close()
```

In [0]:

```

if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
        print("Questions after preprocessed")
        print('='*100)
        reader.fetchone()
        for row in reader:
            print(row)
            print('-'*100)
conn_r.commit()
conn_r.close()

```

Questions after preprocessed

```

=====
=====

```

```

('ef code first defin one mani relationship differ key troubl defin o
ne zero mani relationship entiti ef object model look like use fluent
api object composit pk defin batch id batch detail id use fluent api
object composit pk defin batch detail id compani id map exist databas
tpt basic idea submittedtransact zero mani submittedsplittransact ass
oci navig realli need one way submittedtransact submittedsplittransac
t need dbcontext class onmodelcr overrid map class lazi load occur su
bmittedtransact submittedsplittransact help would much appreci edit t
aken advic made follow chang dbcontext class ad follow onmodelcr over
rid must miss someth get follow except thrown submittedtransact key b
atch id batch detail id zero one mani submittedsplittransact key batc
h detail id compani id rather assum convent creat relationship two ob
ject configur requir sinc obvious wrong',)

```

```

-----
('explan new statement review section c code came accross statement b
lock come accross new oper use way someon explain new call way',)
-----

```

```

-----
('error function notat function solv logic riddl iloczyni list struct
ur list possibl candid solut list possibl coordin matrix wan na choos
one candid compar possibl candid element equal wan na delet coordin c
all function skasuj look like ni knowledg haskel cant see what wron
g',)
-----

```

```

-----
('step plan move one isp anoth one work busi plan switch isp realli s
oon need chang lot inform dns wan wan wifi question guy help mayb peo
pl plan correct chang current isp new one first dns know receiv new i
p isp major chang need take consider exchang server owa vpn two site
link wireless connect km away citrix server vmware exchang domain con
trol link place import server crucial step inform need know avoid dow
ntim busi regard ndavid',)
-----

```

```

-----
('use ef migrat creat databas googl migrat tutori af first run applic
creat databas ef enabl migrat way creat databas migrat rune applic tr
i',)
-----

```

```

-----
('magento unit test problem magento site recent look way check integr
magento site given point unit test jump one method would assum would

```

big job write whole lot test check everyth site work anyon involv uni
t test magento advis follow possibl test whole site custom modul nis
exempl test would amaz given site heavili link databas would nbe poss
ibl fulli test site without disturb databas better way automaticlli c
heck integr magento site say integr realli mean fault site ship payme
nt etc work correct',)

('find network devic without bonjour write mac applic need discov mac
pcs iphon ipad connect wifi network bonjour seem reason choic turn pr
oblem mani type router mine exempl work block bonjour servic need fin
d ip devic tri connect applic specif port determin process run best a
pproach accomplish task without violat app store sandbox',)

('send multipl row mysql databas want send user mysql databas column
user skill time nnow want abl add one row user differ time etc would
code send databas nthen use help schema',)

('insert data mysql php powerpoint event powerpoint present run conti
nu way updat slide present automat data mysql databas websit',)

In [9]:

```
#Taking 1 Million entries to a dataframe.
write_db = 'Processed.db'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags FROM Question
conn_r.commit()
conn_r.close()
```

In [10]:

```
preprocessed_data.head()
```

Out[10]:

	question	tags
0	chang cpu soni vaio pcg grx tri everywher find...	cpu motherboard sony-vaio replacement disassembly
1	display size grayscale qimag qt abl display ima...	c++ qt qt4
2	datagrid selecteditem set back null eventtocon...	mvvm silverlight-4.0
3	filter string collect base listview item resol...	c# winforms string listview collections
4	disabl home button without use type keyguard c...	android android-layout android-manifest androi...

In [11]:

```
print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])
```

```
number of data points in sample : 999999
number of dimensions : 2
```

4. Machine Learning Models

4.1 Converting tags for multilabel problems

X	y1	y2	y3	y4
x1	0	1	1	0
x1	1	0	0	0
x1	0	1	0	0

In [15]:

```
# binary='true' will give a binary vectorizer
vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true')
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```

We will sample the number of tags instead considering all of them (due to limitation of computing power)

In [16]:

```
def tags_to_choose(n):
    t = multilabel_y.sum(axis=0).tolist()[0]
    sorted_tags_i = sorted(range(len(t)), key=lambda i: t[i], reverse=True)
    multilabel_yn=multilabel_y[:,sorted_tags_i[:n]]
    return multilabel_yn


def questions_explained_fn(n):
    multilabel_yn = tags_to_choose(n)
    x= multilabel_yn.sum(axis=1)
    return (np.count_nonzero(x==0))
```

In [17]:

```
questions_explained = []
total_tags=multilabel_y.shape[1]
total_qs=preprocessed_data.shape[0]
for i in range(500, total_tags, 100):
    questions_explained.append(np.round(((total_qs-questions_explained_fn(i))/total
```

In [18]:

```
fig, ax = plt.subplots()
ax.plot(questions_explained)
xlabel = list(500+np.array(range(-50,450,50))*50)
ax.set_xticklabels(xlabel)
plt.xlabel("Number of tags")
plt.ylabel("Number Questions covered partially")
plt.grid()
plt.show()
# you can choose any number of tags based on your computing power, minimum is 50(it
print("with ",5500,"tags we are covering ",questions_explained[50],"% of questions")
```



with 5500 tags we are covering 99.035 % of questions

In [19]:

```
multilabel_yx = tags_to_choose(5500)
print("number of questions that are not covered :", questions_explained_fn(5500),"o
number of questions that are not covered : 9645 out of 999999
```

In [20]:

```
print("Number of tags in sample :", multilabel_y.shape[1])
print("number of tags taken :", multilabel_yx.shape[1],"(",(multilabel_yx.shape[1]/
```

```
Number of tags in sample : 35422
number of tags taken : 5500 ( 15.527073570097679 %)
```

We consider top 15% tags which covers 99% of the questions

4.2 Split the data into test and train (80:20)

In [21]:

```
total_size=preprocessed_data.shape[0]
train_size=int(0.80*total_size)

x_train=preprocessed_data.head(train_size)
x_test=preprocessed_data.tail(total_size - train_size)

y_train = multilabel_yx[0:train_size,:]
y_test = multilabel_yx[train_size:total_size,:]
```

In [22]:

```
print("Number of data points in train data :", y_train.shape)
print("Number of data points in test data :", y_test.shape)
```

Number of data points in train data : (799999, 5500)
 Number of data points in test data : (200000, 5500)

4.3 Featurizing data

In [0]:

```
start = datetime.now()
vectorizer = TfidfVectorizer(min_df=0.00009, max_features=200000, smooth_idf=True,
                             tokenizer = lambda x: x.split(), sublinear_tf=False, n
x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:09:50.460431

In [0]:

```
print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_train.shape)
print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.shape)
```

Dimensions of train data X: (799999, 88244) Y : (799999, 5500)
 Dimensions of test data X: (200000, 88244) Y: (200000, 5500)

In [0]:

```

# https://www.analyticsvidhya.com/blog/2017/08/introduction-to-multi-label-classifi
#https://stats.stackexchange.com/questions/117796/scikit-multi-label-classification
# classifier = LabelPowerSet(GaussianNB())
"""
from skmultilearn.adapt import MLkNN
classifier = MLkNN(k=21)

# train
classifier.fit(x_train_multilabel, y_train)

# predict
predictions = classifier.predict(x_test_multilabel)
print(accuracy_score(y_test, predictions))
print(metrics.f1_score(y_test, predictions, average = 'macro'))
print(metrics.f1_score(y_test, predictions, average = 'micro'))
print(metrics.hamming_loss(y_test, predictions))

"""
# we are getting memory error because the multilearn package
# is trying to convert the data into dense matrix
# -----
#MemoryError                                Traceback (most recent call last)
#<ipython-input-170-f0e7c7f3e0be> in <module>()
#----> classifier.fit(x_train_multilabel, y_train)

```

Out[92]:

```

"\nfrom skmultilearn.adapt import MLkNN\nclassifier = MLkNN(k=21)\n\n# train\nclassifier.fit(x_train_multilabel, y_train)\n\n# predict\npredictions = classifier.predict(x_test_multilabel)\nprint(accuracy_score(y_test, predictions))\nprint(metrics.f1_score(y_test, predictions, average = 'macro'))\nprint(metrics.f1_score(y_test, predictions, average = 'micro'))\nprint(metrics.hamming_loss(y_test, predictions))\n\n"

```

4.4 Applying Logistic Regression with OneVsRest Classifier

In [0]:

```
# this will be taking so much time try not to run it, download the lr_with_equal_weight.pkl
# This takes about 6-7 hours to run.
classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, penalty='l2'))
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)

print("accuracy :",metrics.accuracy_score(y_test,predictions))
print("macro f1 score :",metrics.f1_score(y_test, predictions, average = 'macro'))
print("micro f1 score :",metrics.f1_score(y_test, predictions, average = 'micro'))
print("hamming loss :",metrics.hamming_loss(y_test,predictions))
print("Precision recall report :\n",metrics.classification_report(y_test, predictions))
```

```
accuracy : 0.081965
macro f1 score : 0.0963020140154
micro f1 score : 0.374270748817
hamming loss : 0.00041225090909090907
Precision recall report :
```

	precision	recall	f1-score	support
0	0.62	0.23	0.33	15760
1	0.79	0.43	0.56	14039
2	0.82	0.55	0.66	13446
3	0.76	0.42	0.54	12730
4	0.94	0.76	0.84	11229
5	0.85	0.64	0.73	10561
6	0.70	0.30	0.42	6958
7	0.87	0.61	0.72	6309
8	0.70	0.40	0.50	6032
9	0.78	0.43	0.55	6020
10	0.86	0.62	0.72	5707
11	0.52	0.17	0.25	5723
12	0.55	0.10	0.16	5531

In [0]:

```
from sklearn.externals import joblib
joblib.dump(classifier, 'lr_with_equal_weight.pkl')
```

4.5 Modeling with less data points (0.5M data points) and more weight to title and 500 tags only.

In []:

```
sql_create_table = """CREATE TABLE IF NOT EXISTS QuestionsProcessed (question text
create_database_table("Titlmoreweight.db", sql_create_table)
```

In [24]:

```
pwd
```

Out[24]:

```
'D:\\APPLIEDAI\\Ajay\\stackoverflow'
```


In [30]:

```

# http://www.sqlitetutorial.net/sqlite-delete/
# https://stackoverflow.com/questions/2279706/select-random-row-from-a-sqlite-table
import os
read_db = 'train_no_dup.db'
write_db = 'Titlemoreweight.db'
train_datasize = 400000
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        # for selecting first 0.5M rows
        reader.execute("SELECT Title, Body, Tags From no_dup_train LIMIT 500001;")
        # for selecting random points
        #reader.execute("SELECT Title, Body, Tags From no_dup_train ORDER BY RANDOM")

if os.path.isfile(write_db):
    conn_w = create_connection(write_db)
    if conn_w is not None:
        tables = checkTableExists(conn_w)
        writer = conn_w.cursor()
        if tables != 0:
            writer.execute("DELETE FROM QuestionsProcessed WHERE 1")
            print("Cleared All the rows")

```

Tables in the database:

QuestionsProcessed

Cleared All the rows

4.5.1 Preprocessing of questions

1. Separate Code from Body
2. Remove Special characters from Question title and description (not in code)
3. **Give more weightage to title : Add title three times to the question**

- Remove stop words (Except 'C')
- Remove HTML Tags
- Convert all the characters into small letters
- Use SnowballStemmer to stem the words

In []:

```

#http://www.bernzilla.com/2008/05/13/selecting-a-random-row-from-an-sqlite-table/
start = datetime.now()
preprocessed_data_list=[]
reader.fetchone()
questions_with_code=0
len_pre=0
len_post=0
questions_proccesed = 0
for row in reader:

    is_code = 0

    title, question, tags = row[0], row[1], str(row[2])

    if '<code>' in question:
        questions_with_code+=1
        is_code = 1
    x = len(question)+len(title)
    len_pre+=x

    code = str(re.findall(r'<code>(.*?)</code>', question, flags=re.DOTALL))

    question=re.sub('<code>(.*?)</code>', '', question, flags=re.MULTILINE|re.DOTALL)
    question=stripthtml(question.encode('utf-8'))

    title=title.encode('utf-8')

    # adding title three time to the data to increase its weight
    # add tags string to the training data

    question=str(title)+" "+str(title)+" "+str(title)+" "+question

#     if questions_proccesed<=train_datasize:
#         question=str(title)+" "+str(title)+" "+str(title)+" "+question+" "+str(tags)
#     else:
#         question=str(title)+" "+str(title)+" "+str(title)+" "+question

    question=re.sub(r'[^A-Za-z0-9#+.\-]+', ' ', question)
    words=word_tokenize(str(question.lower()))

    #Removing all single letter and and stopwords from question exceptt for the let
    question=' '.join(str(stemmer.stem(j)) for j in words if j not in stop_words and len(j)>2)

    len_post+=len(question)
    tup = (question,code,tags,x,len(question),is_code)
    questions_proccesed += 1
    writer.execute("insert into QuestionsProcessed(question,code,tags,words_pre,words_post) values(?,?,?,?,?)")
    if (questions_proccesed%100000==0):
        print("number of questions completed=",questions_proccesed)

no_dup_avg_len_pre=(len_pre*1.0)/questions_proccesed
no_dup_avg_len_post=(len_post*1.0)/questions_proccesed

print( "Avg. length of questions(Title+Body) before processing: %d"%no_dup_avg_len_pre)
print( "Avg. length of questions(Title+Body) after processing: %d"%no_dup_avg_len_post)
print( "Percent of questions containing code: %d"%((questions_with_code*100.0)/questions_proccesed))

print("Time taken to run this cell :", datetime.now() - start)

```

In []:

```
# never forget to close the connections or else we will end up with database locks  
conn_r.commit()  
conn_w.commit()  
conn_r.close()  
conn_w.close()
```

__ Sample quesitons after preprocessing of data __

In [0]:

```

if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        reader = conn_r.cursor()
        reader.execute("SELECT question From QuestionsProcessed LIMIT 10")
        print("Questions after preprocessed")
        print('='*100)
        reader.fetchone()
        for row in reader:
            print(row)
            print('-'*100)
conn_r.commit()
conn_r.close()

```

Questions after preprocessed

```

=====
=====

```

```

('dynam datagrid bind silverlight dynam datagrid bind silverlight dyn
am datagrid bind silverlight bind datagrid dynam code wrote code debu
g code block seem bind correct grid come column form come grid column
although necessari bind nthank repli advance..',)
-----
-----

```

```

('java.lang.noclassdeffounderror javax servlet jsp tagext taglibraryv
alid java.lang.noclassdeffounderror javax servlet jsp tagext taglibra
ryvalid java.lang.noclassdeffounderror javax servlet jsp tagext tagli
braryvalid follow guid link instal jstl got follow error tri launch j
sp page java.lang.noclassdeffounderror javax servlet jsp tagext tagli
braryvalid taglib declar instal jstl 1.1 tomcat webapp tri project wo
rk also tri version 1.2 jstl still messag caus solv',)
-----
-----

```

```

('java.sql.sqlexcept microsoft odbc driver manag invalid descriptor i
ndex java.sql.sqlexcept microsoft odbc driver manag invalid descripto
r index java.sql.sqlexcept microsoft odbc driver manag invalid descri
ptor index use follow code display caus solv',)
-----
-----

```

```

('better way updat feed fb php sdk better way updat feed fb php sdk b
etter way updat feed fb php sdk novic facebook api read mani tutori s
till confused.i find post feed api method like correct second way use
curl someth like way better',)
-----
-----

```

```

('btnadd click event open two window record ad btnadd click event ope
n two window record ad btnadd click event open two window record ad o
pen window search.aspx use code hav add button search.aspx nwhen inse
rt record btnadd click event open anoth window nafter insert record c
lose window',)
-----
-----

```

```

('sql inject issu prevent correct form submiss php sql inject issu pr
event correct form submiss php sql inject issu prevent correct form s
ubmiss php check everyth think make sure input field safe type sql in
ject good news safe bad news one tag mess form submiss place even tou
ch life figur exact html use templat file forgiv okay entir php scrip
t get execut see data post none forum field post problem use someth t
itl field none data get post current use print post see submit noth w
ork flawless statement though also mention script work flawless local

```

```
machin use host come across problem state list input test mess',)
```

```
-----
('countabl subaddit lebesgu measur countabl subaddit lebesgu measur c
ountabl subaddit lebesgu measur let lbrace rbrace sequenc set sigma -
algebra mathcal want show left bigcup right leq sum left right counta
bl addit measur defin set sigma algebra mathcal think use monoton pro
perti somewher proof start appreci littl help nthank ad han answer ma
ke follow addit construct given han answer clear bigcup bigcup cap em
ptyset neq left bigcup right left bigcup right sum left right also co
nstruct subset monoton left right leq left right final would sum leq
sum result follow',)
```

```
-----
('hql equival sql queri hql equival sql queri hql equival sql queri h
ql queri replac name class properti name error occur hql error',)
```

```
-----
('undefin symbol architectur i386 objc class skpsmtpmessag referenc e
rror undefin symbol architectur i386 objc class skpsmtpmessag referen
c error undefin symbol architectur i386 objc class skpsmtpmessag refe
renc error import framework send email applic background import frame
work i.e skpsmtpmessag somebody suggest get error collect2 ld return
exit status import framework correct sorc taken framework follow mfma
ilcomposeviewcontrol question lock field updat answer drag drop folde
r project click copi nthat',)
```

__ Saving Preprocessed data to a Database __

In [37]:

```
#Taking 0.5 Million entries to a dataframe.
write_db = 'Titlemoreweight.db'
if os.path.isfile(write_db):
    conn_r = create_connection(write_db)
    if conn_r is not None:
        preprocessed_data = pd.read_sql_query("""SELECT question, Tags FROM Questio
conn_r.commit()
conn_r.close()
```

In []:

```
# never forget to close the conections or else we will end up with database locks
conn_r.commit()
conn_w.commit()
conn_r.close()
conn_w.close()
```

In [39]:

```
preprocessed_data.head()
```

Out[39]:

	question	tags
0	dynam datagrid bind silverlight dynam datagrid...	c# silverlight data-binding
1	dynam datagrid bind silverlight dynam datagrid...	c# silverlight data-binding columns
2	java.lang.noclassdeffounderror javax servlet j...	jsp jstl
3	java.sql.sqlexcept microsoft odbc driver manag...	java jdbc
4	better way updat feed fb php sdk better way up...	facebook api facebook-php-sdk

In [40]:

```
print("number of data points in sample :", preprocessed_data.shape[0])
print("number of dimensions :", preprocessed_data.shape[1])
```

number of data points in sample : 500000
 number of dimensions : 2

Converting string Tags to multilable output variables

In [25]:

```
vectorizer = CountVectorizer(tokenizer = lambda x: x.split(), binary='true')
multilabel_y = vectorizer.fit_transform(preprocessed_data['tags'])
```


Selecting 500 Tags

In [26]:

```
questions_explained = []
total_tags=multilabel_y.shape[1]
total_qs=preprocessed_data.shape[0]
for i in range(500, total_tags, 100):
    questions_explained.append(np.round(((total_qs-questions_explained_fn(i))/total
```

In [27]:

```
fig, ax = plt.subplots()
ax.plot(questions_explained)
xlabel = list(500+np.array(range(-50,450,50))*50)
ax.set_xticklabels(xlabel)
plt.xlabel("Number of tags")
plt.ylabel("Number Questions covered partially")
plt.grid()
plt.show()
# you can choose any number of tags based on your computing power, minimum is 500(i
print("with ",5500,"tags we are covering ",questions_explained[50],"% of questions")
print("with ",500,"tags we are covering ",questions_explained[0],"% of questions")
```



with 5500 tags we are covering 99.035 % of questions
 with 500 tags we are covering 90.025 % of questions

In [28]:

```
# we will be taking 500 tags
multilabel_yx = tags_to_choose(500)
print("number of questions that are not covered :", questions_explained_fn(500),"ou
number of questions that are not covered : 99745 out of 999999
```

In [30]:

```
train_datasize = 400000
x_train=preprocessed_data.head(train_datasize)
x_test=preprocessed_data.tail(preprocessed_data.shape[0] - 400000)
y_train = multilabel_yx[0:train_datasize,:]
y_test = multilabel_yx[train_datasize:preprocessed_data.shape[0],:]
```

In [33]:

```
print("Number of data points in train data :", y_train.shape)
print("Number of data points in test data :", y_test.shape)
```

Number of data points in train data : (400000, 500)
 Number of data points in test data : (599999, 500)

4.5.2 Featurizing data with Tfidf vectorizer

In [52]:

```
start = datetime.now()
vectorizer = TfidfVectorizer(min_df=0.00009, max_features=200000, smooth_idf=True,
                             tokenizer = lambda x: x.split(), sublinear_tf=False, n
x_train_multilabel = vectorizer.fit_transform(x_train['question'])
x_test_multilabel = vectorizer.transform(x_test['question'])
print("Time taken to run this cell :", datetime.now() - start)
```

Time taken to run this cell : 0:04:15.255264

In [53]:

```
print("Dimensions of train data X:",x_train_multilabel.shape, "Y :",y_train.shape)
print("Dimensions of test data X:",x_test_multilabel.shape,"Y:",y_test.shape)
```

Dimensions of train data X: (400000, 95585) Y : (400000, 500)
Dimensions of test data X: (100000, 95585) Y: (100000, 500)

4.5.3 Applying Logistic Regression with OneVsRest Classifier

In [54]:

```

start = datetime.now()
classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, penalty='l2'))
classifier.fit(x_train_multilabel, y_train)
predictions = classifier.predict(x_test_multilabel)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print(metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.23703
Hamming loss  0.00278042
Micro-average quality numbers
Precision: 0.7216, Recall: 0.3259, F1-measure: 0.4490
Macro-average quality numbers
Precision: 0.5492, Recall: 0.2581, F1-measure: 0.3351

```

	precision	recall	f1-score	support
0	0.95	0.64	0.77	5519
1	0.69	0.26	0.38	8190
2	0.81	0.38	0.51	6529
3	0.81	0.43	0.56	3231
4	0.80	0.41	0.54	6430
5	0.82	0.34	0.48	2879
6	0.88	0.49	0.63	5086
7	0.88	0.54	0.67	4533
8	0.62	0.13	0.21	3000
9	0.81	0.52	0.63	2765
10	0.59	0.16	0.26	3051
11	0.70	0.22	0.45	3000

In [56]:

```
import joblib
```

In [57]:

```
joblib.dump(classifier, 'lr_with_more_title_weight.pkl')
```

Out[57]:

```
['lr_with_more_title_weight.pkl']
```

In [58]:

```

start = datetime.now()
classifier_2 = OneVsRestClassifier(LogisticRegression(penalty='l1'), n_jobs=-1)
classifier_2.fit(x_train_multilabel, y_train)
predictions_2 = classifier_2.predict(x_test_multilabel)
print("Accuracy :",metrics.accuracy_score(y_test, predictions_2))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions_2))

precision = precision_score(y_test, predictions_2, average='micro')
recall = recall_score(y_test, predictions_2, average='micro')
f1 = f1_score(y_test, predictions_2, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, rec

precision = precision_score(y_test, predictions_2, average='macro')
recall = recall_score(y_test, predictions_2, average='macro')
f1 = f1_score(y_test, predictions_2, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, rec

print (metrics.classification_report(y_test, predictions_2))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.25123
Hamming loss  0.00270306
Micro-average quality numbers
Precision: 0.7172, Recall: 0.3673, F1-measure: 0.4858
Macro-average quality numbers
Precision: 0.5568, Recall: 0.2950, F1-measure: 0.3709

```

	precision	recall	f1-score	support
0	0.94	0.72	0.82	5519
1	0.70	0.34	0.45	8190
2	0.80	0.42	0.55	6529
3	0.82	0.49	0.61	3231
4	0.80	0.44	0.57	6430
5	0.82	0.38	0.52	2879
6	0.86	0.53	0.66	5086
7	0.87	0.58	0.70	4533
8	0.60	0.14	0.22	3000
9	0.82	0.57	0.67	2765
10	0.60	0.20	0.30	3051
11	0.60	0.20	0.30	3051

5. Assignments

1. Use bag of words upto 4 grams and compute the micro f1 score with Logistic regression(OvR)
2. Perform hyperparam tuning on alpha (or lambda) for Logistic regression to improve the performance using GridSearch
3. Try OneVsRestClassifier with Linear-SVM (SGDClassifier with loss-hinge)

Featurizing data - BOW

In [23]:

```
def bow(train, ds, feature):  
    vectorizer = CountVectorizer(ngram_range=(1, 4), max_features=200000, min_df=0.  
                                tokenizer = lambda x: x.split())  
    vectorizer.fit(train[feature])  
    feature_bow = vectorizer.transform(ds[feature].values)  
    print("Shape of matrix after one hot encoding ", feature_bow.shape)  
    return feature_bow
```

In [34]:

```
question_bow_x_train_ = bow(x_train, x_train, 'question')  
question_bow_x_test_ = bow(x_train, x_test, 'question')
```

Shape of matrix after one hot encoding (400000, 90506)

Shape of matrix after one hot encoding (599999, 90506)

In [37]:

```

start = datetime.now()
classifier = OneVsRestClassifier(SGDClassifier(loss='log', alpha=0.00001, penalty='l2'))
classifier.fit(question_bow_x_train_, y_train)
predictions = classifier.predict(question_bow_x_test_)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print(metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.12518687531145886
Hamming loss  0.004409294015490026
Micro-average quality numbers
Precision: 0.4004, Recall: 0.4446, F1-measure: 0.4213
Macro-average quality numbers
Precision: 0.2921, Recall: 0.3726, F1-measure: 0.3251

```

	precision	recall	f1-score	support
0	0.44	0.42	0.43	47524
1	0.55	0.54	0.55	42561
2	0.63	0.61	0.62	40426
3	0.54	0.50	0.52	37964
4	0.84	0.84	0.84	33571
5	0.72	0.67	0.69	31779
6	0.49	0.46	0.47	20526
7	0.71	0.68	0.69	18810
8	0.52	0.48	0.50	18365
9	0.52	0.51	0.52	17933
10	0.67	0.66	0.66	17263
11	0.33	0.33	0.33	17103

OneVsRestClassifier with Logistic regression

(alpha tuning using Gridsearch)

OneVsRestClassifier with SGDClassifier(penalty=l2, loss=log)==> {Logistic regression}

In [42]:

```

start = datetime.now()
import warnings
warnings.filterwarnings('ignore')

# hp1={'estimator__C':alpha}

alpha=[10**-3,10**-2,10**-1]

cv_scores = []
for i in alpha:
    print(i)
    hp1={'estimator__alpha':[i],
        'estimator__loss':['log'],
        'estimator__penalty':['l2']}
    print(hp1)
    classifier = OneVsRestClassifier(SGDClassifier())

    model11 =GridSearchCV(classifier, hp1,
                          cv=3, scoring='f1_micro', n_jobs=-1)
    print("Gridsearchcv")
    best_model1=model11.fit(question_bow_x_train_, y_train)
    print('fit model')
    Train_model_score=best_model1.score(question_bow_x_train_,
                                         y_train)
# print("best_model1")
    cv_scores.append(Train_model_score.mean())

fscore = [x for x in cv_scores]

# determining best alpha
optimal_alpha21 = alpha[fscore.index(max(fscore))]
print('\n The optimal value of alpha with penalty=l2 and loss= log is %d.' % optimal_alpha21)

# Plots
fig4 = plt.figure( facecolor='c', edgecolor='k')
plt.plot(alpha, fscore,color='green', marker='o', linestyle='dashed',
linewidth=2, markersize=12)

for xy in zip(alpha, np.round(fscore,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Hyper parameter Alpha')
plt.ylabel('F1_Score value ')
plt.show()

print("Time taken to run this cell :", datetime.now() - start)

```

0.001

```
{'estimator__alpha': [0.001], 'estimator__loss': ['log'], 'estimator__penalty': ['l2']}
```

Gridsearchcv

fit model

0.01

```
{'estimator__alpha': [0.01], 'estimator__loss': ['log'], 'estimator__penalty': ['l2']}
```


Gridsearchcv

fit model

0.1

```
{'estimator__alpha': [0.1], 'estimator__loss': ['log'], 'estimator__penalty': ['l2']}\nGridsearchcv\nfit model
```

The optimal value of alpha with penalty=l2 and loss= log is 0.



Time taken to run this cell : 0:56:25.803343

In [43]:

```
print(optimal_alpha21)
```

0.001

In [45]:

```
start = datetime.now()
best_model1 = OneVsRestClassifier(SGDClassifier(loss='log', alpha=optimal_alpha21,
                                              penalty='l2'), n_jobs=-1)
best_model1.fit(question_bow_x_train_, y_train)
```

Out[45]:

```
OneVsRestClassifier(estimator=SGDClassifier(alpha=0.001, average=False,
class_weight=None,
early_stopping=False, eps=
ilon=0.1,
eta0=0.0, fit_intercept=True,
l1_ratio=0.15,
learning_rate='optimal',
loss='log',
max_iter=1000, n_iter_no_
change=5,
n_jobs=None, penalty='l
2',
power_t=0.5, random_state
=None,
shuffle=True, tol=0.001,
validation_fraction=0.1,
verbose=0,
warm_start=False),
n_jobs=-1)
```

In [46]:

```
joblib.dump(best_model1, 'best_model1_LR.pkl')
```

Out[46]:

```
['best_model1_LR.pkl']
```

In [47]:

```
best_model1=joblib.load('best_model1_LR.pkl')
```

In [48]:

```

predictions = best_model1.predict (question_bow_x_test_)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-averasge quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, rec

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, rec

print (metrics.classification_report(y_test, predictions)) #printing classification
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.1896186493644156
Hamming loss  0.0030892151486919144
Micro-averasge quality numbers
Precision: 0.6966, Recall: 0.2558, F1-measure: 0.3742
Macro-average quality numbers
Precision: 0.4868, Recall: 0.1383, F1-measure: 0.2032

```

	precision	recall	f1-score	support
0	0.60	0.15	0.24	47524
1	0.80	0.33	0.47	42561
2	0.84	0.50	0.62	40426
3	0.76	0.39	0.52	37964
4	0.94	0.70	0.80	33571
5	0.87	0.62	0.73	31779
6	0.69	0.21	0.33	20526
7	0.88	0.55	0.68	18810
8	0.73	0.33	0.45	18365
9	0.79	0.37	0.50	17933
10	0.84	0.59	0.70	17263
11	0.52	0.15	0.24	17100

OneVsRestClassifier with Logistic regression(penalty=l1)

In [50]:

```

start = datetime.now()
import warnings
warnings.filterwarnings('ignore')

# hp1={'estimator__C':alpha}

cv_scores = []
for i in alpha:
    print(i)
    hp1={'estimator__alpha':[i],
          'estimator__loss':['log'],
          'estimator__penalty':['l1']}
    print(hp1)
    classifier = OneVsRestClassifier(SGDClassifier())

    model11 =GridSearchCV(classifier, hp1,
                           cv=3, scoring='f1_micro', n_jobs=-1)
    print("Gridsearchcv")
    best_model1=model11.fit(question_bow_x_train_, y_train)
    print('fit model')
    Train_model_score=best_model1.score(question_bow_x_train_,
                                         y_train)
# print("best_model1")
    cv_scores.append(Train_model_score.mean())

fscore = [x for x in cv_scores]

# determining best alpha
optimal_alpha22 = alpha[fscore.index(max(fscore))]
print('\n The optimal value of alpha with penalty=l1 and loss= log is %d.' % optimal_alpha22)

# Plots
fig4 = plt.figure( facecolor='c', edgecolor='k')
plt.plot(alpha, fscore,color='green', marker='o', linestyle='dashed',
         linewidth=2, markersize=12)

for xy in zip(alpha, np.round(fscore,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Hyper parameter Alpha')
plt.ylabel('F1_Score value ')
plt.show()

print("Time taken to run this cell :", datetime.now() - start)

```

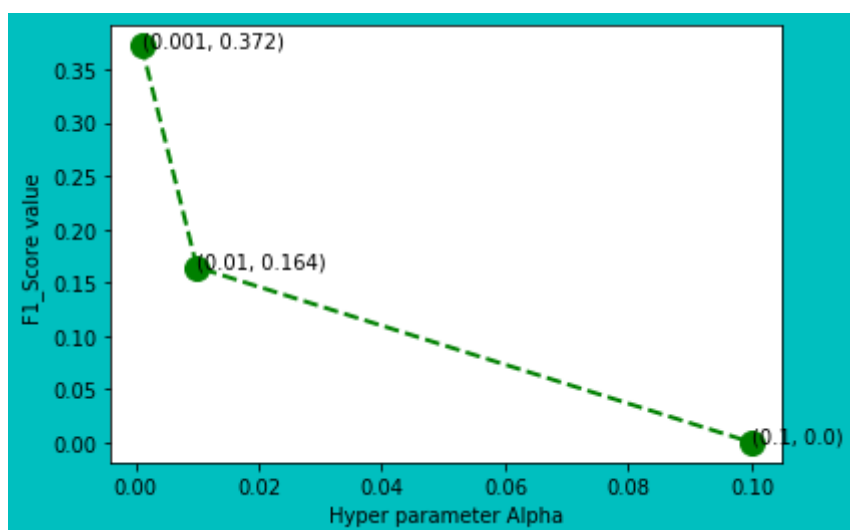
```

0.001
{'estimator__alpha': [0.001], 'estimator__loss': ['log'], 'estimator__penalty': ['l1']}
Gridsearchcv
fit model
0.01
{'estimator__alpha': [0.01], 'estimator__loss': ['log'], 'estimator__penalty': ['l1']}
Gridsearchcv
fit model
0.1
{'estimator__alpha': [0.1], 'estimator__loss': ['log'], 'estimator__penalty': ['l1']}

```

```
_penalty': ['l1']}  
Gridsearchcv  
fit model
```

The optimal value of alpha with penalty=l1 and loss= log is 0.



Time taken to run this cell : 1:45:05.530338

In [51]:

```
optimal_alpha22
```

Out[51]:

0.001

In [52]:

```
start = datetime.now()
best_model2 = OneVsRestClassifier(SGDClassifier(loss='log', alpha=optimal_alpha22,
                                              penalty='l1'), n_jobs=-1)
best_model2.fit(question_bow_x_train_, y_train)
```

Out[52]:

```
OneVsRestClassifier(estimator=SGDClassifier(alpha=0.001, average=False,
class_weight=None,
early_stopping=False, eps
ilon=0.1,
eta0=0.0, fit_intercept=True,
l1_ratio=0.15,
learning_rate='optimal',
loss='log',
max_iter=1000, n_iter_no_
change=5,
n_jobs=None, penalty='l
1',
power_t=0.5, random_state
=None,
shuffle=True, tol=0.001,
validation_fraction=0.1,
verbose=0,
warm_start=False),
n_jobs=-1)
```

In [53]:

```
joblib.dump(best_model2, 'best_model2_LR.pkl')
```

Out[53]:

```
['best_model2_LR.pkl']
```

In [54]:

```
best_model2=joblib.load('best_model2_LR.pkl')
```

Logistic regression with l1 penalty

In [55]:

```

start = datetime.now()
#classifier = OneVsRestClassifier(LogisticRegression(penalty='l1'), n_jobs=-1)
#classifier.fit(x_train_multilabel, y_train)
predictions = best_model2.predict(question_bow_x_test_)
print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))
precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, rec

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, rec

print (metrics.classification_report(y_test, predictions))
print("Time taken to run this cell :", datetime.now() - start)

```

```

Accuracy : 0.18613864356440593
Hamming loss  0.0031306185510309183
Micro-average quality numbers
Precision: 0.6744, Recall: 0.2570, F1-measure: 0.3722
Macro-average quality numbers
Precision: 0.3972, Recall: 0.1512, F1-measure: 0.2012

```

	precision	recall	f1-score	support
0	0.43	0.06	0.10	47524
1	0.79	0.27	0.40	42561
2	0.80	0.45	0.58	40426
3	0.76	0.37	0.49	37964
4	0.90	0.71	0.79	33571
5	0.84	0.64	0.73	31779
6	0.66	0.13	0.21	20526
7	0.83	0.65	0.73	18810
8	0.72	0.32	0.44	18365
9	0.74	0.38	0.50	17933
10	0.77	0.64	0.70	17263
11	0.46	0.00	0.15	17100

OneVsRestClassifier with Linear-SVM (SGDClassifier with loss-hinge)

In [57]:

```

start = datetime.now()
import warnings
warnings.filterwarnings('ignore')

# hp1={'estimator__C':alpha}

cv_scores = []
for i in alpha:
    print(i)
    hp1={'estimator__alpha':[i],
        'estimator__loss':['hinge'],
        'estimator__penalty':['l1']}
    print(hp1)
    classifier = OneVsRestClassifier(SGDClassifier())

    model11 =GridSearchCV(classifier, hp1,
                          cv=3, scoring='f1_micro', n_jobs=-1)
    print("Gridsearchcv")
    best_model1=model11.fit(question_bow_x_train_, y_train)
    print('fit model')
    Train_model_score=best_model1.score(question_bow_x_train_,
                                         y_train)
# print("best_model1")
    cv_scores.append(Train_model_score.mean())

fscore = [x for x in cv_scores]

# determining best alpha
optimal_alpha23 = alpha[fscore.index(max(fscore))]
print('\n The optimal value of alpha with penalty=l1 and loss= log is %d.' % optimal_alpha23)

# Plots
fig4 = plt.figure( facecolor='c', edgecolor='k')
plt.plot(alpha, fscore,color='green', marker='o', linestyle='dashed',
         linewidth=2, markersize=12)

for xy in zip(alpha, np.round(fscore,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Hyper parameter Alpha')
plt.ylabel('F1_Score value ')
plt.show()

print("Time taken to run this cell :", datetime.now() - start)

```

```

0.001
{'estimator__alpha': [0.001], 'estimator__loss': ['hinge'], 'estimator__penalty': ['l1']}
Gridsearchcv
fit model
0.01
{'estimator__alpha': [0.01], 'estimator__loss': ['hinge'], 'estimator__penalty': ['l1']}
Gridsearchcv
fit model
0.1
{'estimator__alpha': [0.1], 'estimator__loss': ['hinge'], 'estimator__penalty': ['l1']}


```

```

r__penalty': ['l1']}
Gridsearchcv
fit model

```

The optimal value of alpha with penalty=l1 and loss= log is 0.



Time taken to run this cell : 1:49:39.366356

OneVsRestClassifier with SGDClassifier for optimal alpha with hinge loss

In [67]:

```
optimal_alpha23
```

Out[67]:

```
0.001
```

In [58]:

```

start = datetime.now()
classifier2 = OneVsRestClassifier(SGDClassifier(loss='hinge',
                                              alpha=optimal_alpha23,
                                              penalty='l1'))
classifier2=classifier2.fit(question_bow_x_train_, y_train)

```

In [59]:

```
joblib.dump(classifier2, 'classifier2.pkl')
```

Out[59]:

```
['classifier2.pkl']
```

In [60]:

```
classifier2=joblib.load('classifier2.pkl')
```

In [61]:

```
predictions = classifier2.predict (question_bow_x_test_)

print("Accuracy :",metrics.accuracy_score(y_test, predictions))
print("Hamming loss ",metrics.hamming_loss(y_test,predictions))

precision = precision_score(y_test, predictions, average='micro')
recall = recall_score(y_test, predictions, average='micro')
f1 = f1_score(y_test, predictions, average='micro')

print("Micro-averasge quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

precision = precision_score(y_test, predictions, average='macro')
recall = recall_score(y_test, predictions, average='macro')
f1 = f1_score(y_test, predictions, average='macro')

print("Macro-average quality numbers")
print("Precision: {:.4f}, Recall: {:.4f}, F1-measure: {:.4f}".format(precision, recall, f1))

print (metrics.classification_report(y_test, predictions)) #printing classification report
print("Time taken to run this cell :", datetime.now() - start)
```

Accuracy : 0.19663032771721287

Hamming loss 0.0030567784279640466

Micro-averasge quality numbers

Precision: 0.6985, Recall: 0.2698, F1-measure: 0.3893

Macro-average quality numbers

Precision: 0.2574, Recall: 0.1644, F1-measure: 0.1858

	precision	recall	f1-score	support
0	0.56	0.03	0.05	47524
1	0.73	0.44	0.55	42561
2	0.72	0.61	0.66	40426
3	0.70	0.46	0.55	37964
4	0.92	0.68	0.78	33571
5	0.80	0.72	0.76	31779
6	0.71	0.11	0.19	20526
7	0.82	0.61	0.70	18810
8	0.69	0.37	0.48	18365
9	0.73	0.44	0.55	17933
10	0.74	0.63	0.68	17263
11	0.00	0.00	0.00	17102

Result

In [71]:

```
from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Sr.No", "MODEL", "FEATURIZATION", "PENALTY", "ALPHA", "LOSS", "MICRO_
x.add_row(["1", 'OneVsRest+SGD Classifier', "Tf-idf", "l1", 0.0001, "log", 0.4858])
x.add_row(["2", 'OneVsRest+SGD(log)=LR', "Bag-of-words", "l2", 0.001, "log", 0.3742])
x.add_row(["3", 'OneVsRest+SGD(log)=LR', "Bag-of-words", "l1", 0.001, "log", 0.3722])
x.add_row(["4", 'OneVsRest+SGD Classifier', "Bag-of-words", "l1", 0.001, "Hinge", 0.3893])

print(x)
```

Sr.No	MODEL	FEATURIZATION	PENALTY	ALPHA	LOSS	MICRO_F1_SCORE
1	OneVsRest+SGD Classifier	Tf-idf	l1	0.0001	log	0.4858
2	OneVsRest+SGD(log)=LR	Bag-of-words	l2	0.001	log	0.3742
3	OneVsRest+SGD(log)=LR	Bag-of-words	l1	0.001	log	0.3722
4	OneVsRest+SGD Classifier	Bag-of-words	l1	0.001	Hinge	0.3893

Observation

1. Best Micro F1-score is obtained from TFIDF vectorizer
2. Although we used BOW with ngram=(1,4), but still TFIDF vectorizer has perofrmed better